

Anuj Chaudhary

AI Engineer | Data Scientist | IIT Madras BS (Data Science) | Ex-British Airways | Building Production-Ready ML, Generative AI & LLM Applications

Contact

+91-7017707607

beinganujchaudhary@gmail.com

BeingAnujChaudhary  | LinkedIn

Education

Bachelor of Arts (Majors in Psychology) | IGNOU

2018 – 2022

BS Degree in Data Science and Applications | IIT Madras

2026 – pursuing

Certifications

 Mathematics for Machine Learning by Imperial College London (Coursera)

 Microsoft Power BI Data Analyst Certificate by Microsoft

 Machine Learning Specialization by DeepLearning.AI (Andrew Ng)

 Google IT Support Professional Certificate by Google (Coursera)

 Atlassian IT Service Management (ITSM) Professional Certificate

 Atlassian Agile Project Management Professional Certificate

 ServiceNow IT Leadership Professional Certificate

 AWS Certified Cloud Practitioner by AWS Training & Certification

Objective

Data Science student (IIT Madras) with psychology background and L2 incident management experience, building end-to-end analytics solutions for customer retention and marketing optimization.

Proficient in Python, statistical testing, and interpretable ML (XGBoost, SHAP), with a focus on translating data insights into actionable business recommendations.

Experience

JUNE 2026 – Present

Data Scientist Intern | Tera Global Enterprises

-  Analyzed and processed structured and unstructured data to extract actionable insights.
-  Built and maintained data pipelines for data collection, preprocessing, and transformation.
-  Performed data cleaning, feature engineering, and exploratory data analysis (EDA) to improve data quality.
-  Collaborated with cross-functional teams to support data-driven decision-making and model development.

JUNE 2024 – June 2026

Service Advisor | Call BA – Subsidiary of British Airways

Provided L2 technical and functional support for ticketing, reissues, and fare rule validation using Amadeus GDS.

-  Managed end-to-end PNR lifecycle in Amadeus, including creation, fare pricing, ticketing, reissues, and refunds
-  Analysed PNR data to identify booking inconsistencies and system processing issues across reservation workflows

NOVEMBER 2022 – MAY 2024

Senior Technical Analyst | Inergizer IT Services Private Limited

-  Resolved Level 2 production incidents by analysing logs, application workflows, and user impact
-  Conducted root cause analysis (RCA) to identify recurring system defects and reduce incident volume

Projects

LLM Concepts RAG System | Elasticsearch, Streamlit, Docker, Groq and OpenAI

- ✚ Advanced Retrieval-Augmented Generation for Machine Learning Education
- ✚ An end-to-end RAG application that answers ML questions using hybrid search (BM25 + Vector) and cross-encoder reranking.
- ✚ It features dual-LLM fallback (Groq & OpenAI), dynamic Mermaid diagram generation for visual learning, and a fully Dockerized architecture with PostgreSQL & Grafana for monitoring interactions.

SaarthiEO — Disaster Detection from Satellite & Drone Imagery | Python, PyTorch, Gradio, Hugging Face | [GitHub Link]

- ✚ Fine-tuned EfficientNet-B3 on a 4-class aerial disaster dataset achieving 97.38% test accuracy (Fire, Flood, Collapsed Building, Traffic Incident).
- ✚ Built an end-to-end Grad-CAM explainability pipeline hooking the last convolutional block to produce interpretable attention heatmaps.
- ✚ Deployed a production-grade 3-tab Gradio web app (single predict, Grad-CAM, batch CSV export) to Hugging Face pages with Git LFS model storage.

aetherRead — AI-powered Offline-First PDF Research Workspace | Python, Next.js, TypeScript, Python, Ollama, ChromaDB | [GitHub Link]

- ✚ Designed and engineered a cross-platform, production-grade PDF reader prioritizing privacy and offline-first performance, ensuring zero data dependency on external cloud services.
- ✚ Implemented a local-first storage architecture utilizing IndexedDB (Dexie.js) for high-performance management of PDFs, reading state, and user-generated annotations.
- ✚ Architecting a local RAG pipeline using Ollama (llama3.2:3b) and ChromaDB to enable intelligent document querying, semantic search, and AI-assisted research with verifiable citations.
- ✚ Engineered a "Comfort Engine" featuring six customizable reading themes and page-level annotation capabilities categorized for research workflows (Important, Definition, Question, Quote).

Customer Churn Prediction & Retention Prioritization | Python, XGBoost, SHAP, Power BI | [GitHub Link]

- ✚ Focus: Churn prediction, A/B testing, behavioral segmentation for proactive customer retention
- ✚ Built churn prediction model (XGBoost, 87% AUC) using transactional and behavioral features from 10K+ credit card customers
- ✚ Applied SHAP values to identify top 5 churn drivers (e.g., declining transaction frequency, product disengagement); created risk-based prioritization tiers (High/Medium/Low) for targeted interventions
- ✚ Developed Power BI dashboard enabling sales teams to filter at-risk accounts by value tier, supporting proactive retention workflows
- ✚ Designed reproducible pipeline with synthetic data generator and automated evaluation; documented limitations and monitoring considerations for production readiness

Key Skills

- ✚ **Programming:** Python (NumPy, Pandas, Scikit-learn, XGBoost, SHAP), SQL
- ✚ **Data analysis & visualisation:** Power BI (DAX, interactive dashboards), Matplotlib, exploratory data analysis (EDA), feature engineering
- ✚ **Statistics:** Hypothesis testing, A/B testing, regression, classification, clustering, power analysis, RFM analysis
- ✚ **Tools:** Git, GitHub, Jupyter Notebook, Amadeus GDS, Salesforce, Zendesk, Jira (Atlassian)